

Data-Driven Identification of Elite Soccer Talent via Supervised Learning

Alexander Shipley (alshipley39@gmail.com)

Maxwell Dalrymple (maxwelldalrymple4@gmail.com)

Abstract

Every year a list of the top 100 international football players is released to the public [1]. These rankings are provided by The Guardian newspaper, with the rankings decided via voting rounds conducted by retired players, current coaches and technical staff, and an array of media personnel across the globe. Judges choose their own criterion when selecting their top candidates [1], although it is reasonable to assume votes are heavily based on individual player statistics as well as team performance, where both club and national team performances are weighed if applicable. Prestigious awards such as the Ballon d’Or are presented to players who have demonstrated excellence and dominance in comparison to their peers in the football domain [2]. These award ceremonies are based solely on a single season’s performance, and not a player’s career outright. Due to the nature of football scheduling, awards are almost entirely retroactive [2]. The question our research group set out to answer was to what extent does a player’s performance over their playing career impact their chances of claiming a spot on the top 100 rankings for the current year? Thus, we aimed to build a machine learning algorithm that used data from multiple past seasons to accurately predict whether a player will appear on the top 100 rankings.

Introduction

As long time football fans, we have always been curious how The Guardian determines the top 100 rankings for any given year. We hypothesized that it may be determined by a plethora of factors such as team notoriety, adjacent awards presented to teammates, a player’s public status and popularity, fan viewership and support, player salary, and performances throughout the previous season. After obtaining a Masters of Data Analytics degree [3], our paper’s lead, Alexander, was interested in applying machine learning methodologies to investigate whether we could systematically predict complex multi-variable trends associated with the football ecosystem. We decided to work on an algorithm related to the top 100 ranked football players provided by The Guardian, as we thought it would be interesting to pre-determine our own listings before publication.

For this analysis, data from club football matches spanning from the 2017/2018 (2017/18) season up to and including the 2024/25 season was sourced. This data consists of matches from the Premier League, Ligue 1, La Liga, Bundesliga, and Serie A, also known as the “Big Five” [4]. The primary objec-

tive was to create a predictive model that correctly partitions at least 50 of the top 100 players appearing in a top 100 list for 2025. In other words, our model produces at least 50 true positives (TPs). Partitioning a player correctly takes precedence over the closeness of their predicted rank to their true ranking. National team data was not used, as national team events are not held annually and there exists significant gaps in the skill level of national teams. Of the ten largest reported victories in football history, seven of the matches were contested between national teams [5]. Thus, players registered to national teams would have the opportunity for stat padding against weaker teams, which would in our opinion falsely exaggerate their abilities.

Five distinct predictive models were selected for empirical evaluation. These models consisted of logistic regression; support vector machines (SVMs) with radial basis function (rbf), sigmoid, and polynomial kernels; and a random forest. Following training on an identical dataset comprising historical player match records, each model generates calibrated probability estimates of each player in the test set ranking within the top 100. The player owning the highest predicted probability has the highest likelihood of being in the top 100 rankings. We hypothesized that there is a strong correlation between our models’ predictive strength and an official ordering of the top 100 rankings. Therefore, the player with the highest predicted percentage of being in the top 100 has the highest true ranking (#1 seed).

Data

The data used in this analysis was sourced from FBref.com, a database maintaining historical football matches and player statistics [6]. An existing data harvester written in Python assisted in obtaining the fixture and player data from matches in the Big Five [7]. Fixture data served as a secondary form of verification and contained the match date, the participating teams, and the final score. Player data mainly consisted of in-game related statistics such as minutes played, offensive and defensive stats, penalties, passing, dribbling, discipline (yellow and red cards), and expected stats. Also included was a player’s name, age, shirt number, nationality, and position(s). Goalkeepers (keepers) had additional stats such as shots faced, throws attempted, and goal kicks, among others [7].

With the player data in hand, The Guardian’s annual top 100 reports identified the top players for our training data seasons [8–14]. The annual Guardian reports are mostly published in December, whereas the seasons of the Big Five span from

August to May [4], leading to potential inconsistencies in our dataset. As a methodological compromise, the annual report was reclassified as a seasonal report, such that The Guardian's 2018 report pertains exclusively to matches from the 2017/18 season. This approach, though not empirically factual, allowed us to utilize the specific data format we harvested. This constraint also introduced a misalignment that created potential vulnerabilities in ranking integrity. Players exhibiting poor performance metrics during the concluding months (January - May) of a season could subsequently demonstrate exceptional performance during the initial months (August - December) of the next season, thereby exploiting the seasonal boundary to artificially elevate their ranking within the top 100 classification system. However, with by-season reports scarce, this was the best course of action. At the time of this paper's conception, The Guardian had yet to publish the 2025 report, but the report of the Spanish Newspaper Marca detailing the top 100 players in the 2024/25 season was used [15]. The accuracy of this by-season report would outweigh The Guardian's yearly report with zero risk of seasonal crossover, though the training set rankings consisted of data distributed by The Guardian, which is noteworthy.

We utilized a multitude of Python scripts to match player data to published rankings. Although we have prior experience in other readily-available statistical modeling languages such as R, most of our experience in machine learning and data manipulation is derived from Python implementations. Thus, we decided to do all our automation via Python. Rankings were transformed into a binary notation (0 - not in top 100, 1 - in top 100). This transformation was necessary as our stated goal was classification, not exact ranking placement.

Having executed the majority of our machine learning experience with Python language in the Google Colaboratory (colab) environment, model training and testing were performed on this platform. Preprocessing on the now complete dataset included removing headers and attributes we believed would not impact model predictions. These included indexes, jersey numbers, player nationalities, and unique match IDs. The age attribute, originally encoded in a 'year-day' format, was decomposed into one quantitative attribute: Age Years. The same was done to the `source_file` field (`20##-20##_leagueName`) to extract a season attribute.

Due to the distinct positional attributes of keepers, which prevents direct comparisons with outfielders, we elected to segregate the keepers from the outfielders. We made the executive decision to only test models on the outfielder data due to their persistent dominance in the top 100 rankings. For outfielders, all keeper stats were dropped. Special care had to be taken with the position attribute, as players can be forced to switch positions mid-game. In instances where this occurred, multiple positions encoded a single attribute. Thus, we separated the main position attribute into 14 boolean position columns, each representing if that position was played in the match or not.

We required exactly one instance of our players' stats for each individual season. Thus, our dataset required aggregation of seasonal player stats by either summing or averaging specific attributes. In-game attributes disclosing pass completion percentage and position booleans were summed. The aforementioned along with age in years were aggregated by seasonal averages. Upon completion, the dataset was split into training and testing sets using the season descriptor. The `StandardScaler` function was chosen to normalize our dataset. Standardization is deployed when the range of attributes in a dataset varies. Due to the plethora of attributes and the nature of sports, this was the case for our dataset. For example, in a given season, minutes played can range from 0-1000, while goals are realistically bounded by a range of 0-100. `StandardScaler` normalizes all values around 0 to mitigate attribute dominance [16]. It is noteworthy to add that the testing set was not included in the creation of the standardized object.

Methods

To evaluate our hypothesis, our baseline models were trained and rigorously tested. All model parameters used default values other than a rotating `class_weight` designation to gauge a reasonable standard. The `class_weight` parameter can drastically change the characteristics of a given model, so using all the potential `class_weight` settings was imperative for cross model comparison. Modification of the `class_weight` parameter determines how the model balances classes in the training phase, specifically used to embolden classes with fewer instances than the mean [17].

Unfortunately, our test data did not house 100 true instances of our top 100 player class. First, due to segregating the outfielders from the keepers and exclusively using the outfielder dataset for both training and testing, eight keepers identifying as top 100 players were left out. Next, some of the top 100 players have left the Big Five and have moved to leagues abroad such as Major League Soccer (MLS) or the Saudi Pro League (SPL) [18]. This left only 77 of the top 100 players present in our test dataset.

For baseline testing, the logistic regression and each SVM model's `class_weight` parameter was set to `None` or `balanced`. The random forest's `class_weight` was modified thrice with values `None`, `balanced`, and `balanced subsample`: a `class_weight` exclusive to certain tree models relying on bootstrap sampling [19, 20]. This provided 11 models to inspect.

After baseline testing, each model underwent hyperparameter tuning in hopes of increasing predictive capabilities. `RandomizedSearchCV` (RSCV) was utilized as our hypertuning algorithm, specializing in fixed number random parameter combinations [21]. This paper delves deeper into our rationale behind choosing the RSCV algorithm in later sections. Due to our limited familiarity with numerous model parameters, Generative AI (Gen AI) was utilized to produce valid parameter grids [22].

Parameter	Explanation
penalty	Specifies the regularization technique to control attribute coefficients [23].
C	Controls the strength of the penalty. The smaller the value, the stronger the penalty, and vice versa [24].
solver	Optimization Algorithm. Chosen based on the penalty [24].
max_iter	Maximum number of training steps the solver can run to achieve convergence [24].
l1_ratio	Controls the mix between strengths of the L1 and L2 penalties [24].
gamma	Controls the influence of any single data point [25].
n_estimators	Number of trees in the forest. More trees equals a better performance, but slower model generation [26].
max_depth	The maximum depth of each tree. A higher depth results in more information being recorded about the data, but also increases the risk of overfitting [26].
min_samples_split	"The minimum number of samples required to split an internal node" [26, para. 19]. More samples results in underfitting [26].
min_samples_leaf	"The minimum number of samples required to be at a leaf node" [26, para. 21]. More samples results in underfitting [26].
max_features	The maximum "number of features to consider when looking for the best split" [26, para. 23]. More features results in overfitting [26].
bootstrap	Boolean representing if each tree is built via random subsampling with replacement or using the full dataset [19, 27].

Table 1: Hypertuning Parameter Descriptions

Parameter	Values
penalty	['l1', 'l2', 'elasticnet', None]
C	[0.1, 0.31, 0.52, 0.73, 0.94, 1.15, 1.36, 1.57, 1.78, 2.0]
solver	['lbfgs', 'liblinear', 'saga']
max_iter	[500, 1000, 2500, 5000]
l1_ratio	[0.0, 0.25, 0.50, 0.75, 1.00]
class_weight	['balanced', None]

Table 2: Logistic Regression Parameter Grid

Parameter	Values
C	[0.01, 0.1, 1, 10, 100]
gamma	['scale', 'auto', 0.01, 0.1]
class_weight	['balanced', None]
degree (polynomial kernel)	[2, 3]

Table 3: Support Vector Machine (SVM) Parameter Grid

Parameter	Values
n_estimators	[100, 200, 300, 400, 500, 600, 700, 800, 900, 1000]
max_depth	[None, 5, 10, 20, 30, 50]
min_samples_split	[2, 5, 10, 20]
min_samples_leaf	[1, 2, 4, 10]
max_features	['sqrt', 'log2', None]
bootstrap	[True, False]
class_weight	[None, 'balanced', 'balanced_subsample']

Table 4: Random Forest Parameter Grid

The parameter space of each SVM kernel never exceeded 100 combinations, giving us the freedom to test them all. For consistency, 100 random combinations were chosen for the logistic regression and random forest models. This number saved resource time and kept a similar number of iterations for all models.

Within the RSCV object is a scoring parameter having multiple options to assist in grading the performance of the models. None of the options utilized true positives, our sole indicator for cross model evaluation [21]. Consequently, we developed a custom scoring function; an inaugural feat for both of us despite our prior experience with machine learning algorithms. Our custom function accepted a model, the test data, and the top 100 class values. Each player's predicted probability is outputted and then sorted in descending order. Taking the top 100 predicted probabilities and the true top 100 list, we evaluate the predictions and return the TP percentage in decimal notation. For instance, 57 players being correctly partitioned returns 0.57 (57%).

For cross validation, GroupKFold (GKF) was used. Cross validation is used to gain a more realistic picture of model performance by splitting the training data into groups of multiple training sets and a singular validation set, using each group to train and test the model for further tuning [28]. After reviewing a list of 15 different cross validation methods [29], GKF was deemed the most suitable. In this form of cross validation, a 'group' parameter decides the training vs validation set for a given split, making this algorithm a strong candidate for groupable datasets [29]. GKF complimented our dataset, as our top 100 players varied per season, and the season attribute tracked when matches occurred. Our training set aggregated seven seasons, with six of the seasons acting as training data and one as a validation set in each split. By

making use of the season attribute, we ensured that there was no overlap and that each season’s data was kept in its own fold.

After cross validation was complete, the resulting tuned models were trained and tested using the same process as the baseline models. We then evaluated the tuned models’ performances against their respective baselines.

Results

When running through the 11 baseline tests, we took notice of the number of true positives along with a modified true positive percentage (MTP) after considering that the test data did not include the full list of the top 100 players. The MTP equates to the number of TPs over 77 instead of 100. It is important to note that the MTP percentage is not a true accurate representation of what would occur if all 100 players were present. It is almost certain that the percentage would be slightly higher than originally obtained, but no guarantees as to the extent.

Model	TP	MTP
LogisticRegression (class_weight = ‘balanced’)	47%	61.04%
LogisticRegression (class_weight = None)	47%	61.04%
SVM (kernel = ‘rbf’, class_weight = ‘balanced’)	36%	46.75%
SVM (kernel = ‘rbf’, class_weight = None)	51%	66.23%
SVM (kernel = ‘sigmoid’, class_weight = ‘balanced’)	9%	11.69%
SVM (kernel = ‘sigmoid’, class_weight = None)	29%	37.66%
SVM (kernel = ‘poly’, class_weight = ‘balanced’)	42%	54.55%
SVM (kernel = ‘poly’, class_weight = None)	45%	58.44%
RandomForestClassifier (class_weight = ‘balanced’)	42%	54.55%
RandomForestClassifier (class_weight = None)	51%	66.23%
RandomForestClassifier (class_weight = ‘balanced_subsample’)	40%	51.95%

Table 5: Test Model Percentages

Reviewing Table 5 above and Figure 1 in Appendix A, the performance of the sigmoid kernels were extremely poor, with the balanced rbf and the random forest models having the strongest predictive power. Accounting for only 77 true top 100 instances in our test data, all models save the balanced rbf and the sigmoid kernels can be assumed to have met our goal with the MTP percentage exceeding 50%. To reiterate, the MTP percentage is not the gold standard, but with most TP percentages fairly close to 50% and leaping from 77 to 100

players leading to an almost certain increase in percentage, our assumption is reasonable.

The results of running our RSCV with GKF cross validation on the training data generated the best hyperparameter combinations. The testing data was excluded as it was imperative to test the model on unseen data. Thus, once the tuned parameters were obtained, we could generate new models, re-train them, and test them using the excluded testing data. We obtained the hypertuned percentages after secondary training and testing on the original standardized dataset. The baseline percentages in Table 6 below and Figure 2 in Appendix A were generated by matching the tuned class_weight parameter with that of the corresponding baseline model.

Model	Baseline	Hypertuned
LogisticRegression (C = 0.7333, l1_ratio = 1.0, max_iter = 500, penalty = ‘l1’, solver = ‘saga’, class_weight = None)	TP: 47%, MTP: 61.04%	TP: 49%, MTP: 63.64%
SVM (kernel = ‘rbf’, class_weight = None, C = 1, gamma = 0.01)	TP: 51%, MTP: 66.24%	TP: 52%, MTP: 67.53%
SVM (kernel = ‘sigmoid’, class_weight = None, C = 1, gamma = 0.01)	TP: 29%, MTP: 37.66%	TP: 13%, MTP: 16.88%
SVM (kernel = ‘poly’, class_weight = None, C = 0.1, gamma = ‘scale’, degree = 3)	TP: 45%, MTP: 58.44%	TP: 48%, MTP: 62.34%
RandomForestClassifier (class_weight = ‘balanced_subsample’, max_depth = 20, min_samples_leaf = 4, min_samples_split = 10, n_estimators = 700)	TP: 40%, MTP: 51.95%	TP: 46%, MTP: 59.74%

Table 6: Full Percentages

The sigmoid kernel was the sole model that failed to improve after tuning, but this was expected as our tests established that a sigmoid kernel fits our dataset poorly. The greatest improvement came from the random forest model, where we noticed the tuned model preferred the balanced_subsample class weight rather than None, which had the highest TP percentage of all the random forest models.

Every model has its own unique attribute hierarchy. We opted not to perform feature selection, instead removing attributes we deemed unimportant during the data sanitation stage. To better understand model performance, we wanted deeper insight into the attribute landscape. By evaluating which vari-

ables had the greatest impact on a prediction, our hope is that future models can be improved upon. To gain this understanding, each model had its own procedure to gauge feature importance. For the logistic regression model, we let the coefficients of each attribute determine feature significance. Coefficient magnitude is the simplest method to derive attribute hierarchy in logistic regression [30]. These values can be either positive or negative. A positive coefficient corresponds to increased log odds (and therefore increased odds) of a player belonging in the top 100 assuming a one-unit increase in the inspected feature's value only. A negative coefficient has the opposite effect [30]. The random forest model includes a built-in feature_importances_ parameter, which we selected for use [31]. Regarding the SVM models, coefficients could not be used, as they are only available for linear kernels. Thus, we used Permutation Importance, a ranking method that compares model prediction capabilities before and after values of a given feature are randomly shuffled [32].

For logistic regression, on further analysis of the coefficients, a one unit increase in the features passes completed, pass completion percentage, or goals results in the greatest odds increase of being a top 100 player. A one unit increase in minutes resulted in by far the greatest decrease in the placement odds.

Each model had varying importances for the attributes, but one pattern emerged between nearly all of them: the most important feature was primarily an offensive-minded action. For the random forest, rbf kernel, and polynomial models, the most important features were goal contributing actions, goals, and goals again, respectively. For the sigmoid kernel, the most important feature was the Center Back (CB) position, which strays from the pattern of the other models.

In regards to the positional variables, most of them had very little importance in nearly all the models, but offensive positions tended to rank higher than defensive positions. In the logistic regression model, the offensive positions Forward (FW) and Left Wing (LW) retained the highest positive coefficients of all positional variables. The defensive positions Left Back (LB), Right Back (RB), Defensive Midfielder (DM), and Defender (DF) all held negative coefficients. Surprisingly, the Right Wing (RW) and Attacking Midfielder (AM) positions possessed a negative coefficient, while the CB and Wing Back (WB) positions held positive coefficients. For the feature rankings derived from the feature_importances_ parameter and the Permutation Importance method, offensive positions also primarily held more importance than defensive positions. FW was one of the top two most important positions in all models except the polynomial kernel, and LW was also highly important in most models. Defensive positions did hold some importance, but were deemed less important in the overall picture.

We were also interested in our models' ranking capabilities compared to the official 2025 top 100 players produced by the

Spanish newspaper Marca. We had originally hypothesized a strong correlation between the predicted probability intensities and publicly released rankings. Using the probabilities' strengths, we were able to create ranking lists for each model to further investigate our hypothesis, shown in Figures 3-5 in Appendix A.

After surface level analysis, there does not seem to be a strong correlation present between our prediction strengths and true rankings, as less than 1% of all possible placements match Marca's. A correlation between prediction strength and confirmed TPs does seem to be apparent, demonstrating each model's confidence in its own predictive capabilities.

Discussion

As previously stated, our goal was to evaluate if a model trained on existing footballer data could identify 50 or more players in the top 100 released by Marca. After hyperparameter tuning, the rbf kernel was the sole model to succeed. The logistic regression, polynomial kernel, and random forest models fell shy, with the sigmoid kernel utterly failing. However, when reviewing the MTP rate, it can be estimated that all hypertuned models succeeded in reaching our goal save the sigmoid kernel.

One intriguing factor to discuss is how after tuning, the hypertuned sigmoid and random forest models actually produced fewer TPs than the baseline models. However, it is believed that these baseline results were caused by lucky splitting, and that performing GKF cross validation helped our models better identify patterns in the data. Thus, perhaps if we incorporated the test season into our training data and utilized the future 2026 rankings as new test data, the hyperparameters found from the search would have been better suited for our dataset. Alternatively, using additional training seasons prior to 2017/18 may have added more context to improve models. Furthermore, the random forest model grid search was limited to only 100 out of a possible 17,280 combinations, realistically missing a fully optimized parameter set.

Limiting the grid search was a necessary trade off due to time and computation constraints. The polynomial search took approximately eight hours for completion with only 80 combinations. Given unlimited resources, more combinations would be attempted, potentially yielding better results. GridSearchCV was our original choice for tuning, but as this would have resulted in attempting every possible combination of parameters, we opted for RSCV [33].

When reviewing each hypertuned model's attribute hierarchy, all models excluding the sigmoid kernel favoured offensive attributes. However, this model performed poorly in all aspects, possibly due to the chosen defensive CB position as the leading attribute. In addition, when cross-referencing Marca's chosen players with FBref, the majority of the top players are offensive-minded. Of Marca's 92 outfield players, only 15 are

pure defenders, with another 7 primarily playing as fullbacks or defensive midfielders. At least 40 of the players play a purely forward or attacking midfielder role, with the number likely to be larger due to FBref mixing positions in player bios (55 players are listed as FW/MF for example). Thus, offensive statistics being prioritized aligns with Marca's rankings. In fact, in the 69-year history of the Ballon d'Or, the winners have also been primarily attack-minded, where 27 times the trophy was awarded to a striker, 23 times a winger, and 15 times a midfielder, but only thrice to a defender, and once a keeper [34]. There is a clear bias in favor of attack-minded players being deemed the best players in the world.

Regarding bias, we have identified another potential source that could contribute to skewed results. Some players have been professional footballers for the full eight seasons while others have only just begun their careers. Lamine Yamal is an example of the latter. He only started playing professional football a few years ago [35] and placed first in Marca's rankings [15].

Our decision to either sum or average specific attributes in each player's seasonal data may have also added inherent bias. We believed countable in-game attributes were to be summed, while the rest averaged. This choice, although applied consistently, could still have led to an undetectable reduction in prediction strength.

One last factor for consideration is that the top 100 attribute in our test and training data were not from the same source, as The Guardian's 2025 report had not yet been published as of the time we initialized this research. If we had ensured the same source was used throughout, perhaps each model would have performed better.

In retrospect, if this project were to be attempted again, strategic improvements could be implemented in hopes of boosting model performance. One alteration would be to attempt to predict the true ordering of the rankings instead of using binary classification. This would change our methodology to a Machine-learned Ranking (MLR) task [36]. Regardless of a methodology shift, different models such as neural networks (e.g. MLP) or decision trees could be leveraged. Seeing how our sigmoid kernel yielded poor results, perhaps another model we did not test would be better suited for our dataset. Implementing feature selection could also help optimize models by statistically sorting features by importance, allowing for much more refined predictions. Additionally, we must ensure that the rankings in the training and testing data stem from a common source. When the project commenced, limitations constrained us to use different sources, but now that The Guardian's 2025 report has been published [37], sources can be homogenous. Finally, as many of the removed top 100 players now reside in the MLS and the SPL [18], perhaps it would have been wise to obtain data from these leagues as well.

Conclusion

This paper spawned by admiring The Guardian's yearly evaluation of the top football players in the world, and questioning if annual rankings could be reproduced utilizing machine learning methodologies. Armed with the annual reports of The Guardian and Marca for verification, we had a goal to create a model trained on past match data, with said model achieving a true positive rate of at least 50% when tested on data from the most recently completed season (as of May 2025). This goal was achieved by multiple models, where the rbf kernel and logistic regression model performed the strongest, followed by the polynomial kernel and random forest. In stark contrast, the sigmoid kernel performed poorly in all aspects.

Regarding attribute importance, successful models prioritized offensive statistics higher than defensive, aligning with Marca's chosen players and the previous history of Ballon d'Or winners. Positional variables were prioritized much less than in-game statistics.

Our hypothesis that there was a strong correlation between prediction strengths and Marca's rankings turned out to be false, with less than 1% of players correctly placed. However, there appears to be a correlation between the prediction strength and the TPs, as the stronger the prediction strength, the more likely the player identifies as a TP.

This paper leaves many opportunities to be improved upon in the foreseeable future. Ensuring that the rankings in the training and testing data originate from the same source is imperative, but other options exist too. Whether it be by implementing feature selection, attempting different models, or by changing the type of task from classification to MLR, substantial improvements can be resulted.

References

- [1] M. Christenson, "How the Guardian ranked the 100 best male footballers in the world 2024," *The Guardian*, Dec. 16, 2024. [Online]. Available: <https://www.theguardian.com/football/2024/dec/16/how-the-guardian-ranked-the-100-best-male-footballers-in-the-world-2024>.
- [2] beIN SPORTS, "How the 2025 Ballon d'Or Is Decided: Rules, Voting, and Who Makes Up the Jury," *beIN SPORTS*, Sep. 22, 2025. [Online]. Available: <https://www.beinsports.com/en-us/soccer/articles/how-the-2025-ballon-d-or-is-decided-rules-voting-and-who-makes-up-the-jury-2025-09-22>.
- [3] Western University, *Master of Data Analytics*, <https://www.uwo.ca/mda/index.html>.
- [4] A. Noronha, "When is soccer season? MLS, Youth, Europe start and end date explained," *Goal.com*, Mar. 9, 2025. [Online]. Available: <https://www>

- .goal.com/en-us/lists/season-in-soccer/blt12f2705ce4ea7fb8.
- [5] B. Lane, "Biggest Wins in Soccer History: Ranking The Top Ten," *Sports Illustrated*, Jun. 22, 2025. [Online]. Available: <https://www.si.com/soccer/biggest-wins-in-soccer-history-ranking-top-10>.
 - [6] FBref, *All About FBref.com*, <https://fbref.com/en/about/>.
 - [7] A. Corren, *FBREF Football Player Data Scraper*, https://github.com/adamcorren/fbref_football_player_data_scraper, Jul. 4, 2023.
 - [8] The Guardian, "The 100 best male footballers in the world 2018," *The Guardian*, Dec. 20, 2018. [Online]. Available: <https://www.theguardian.com/football/ng-interactive/2018/dec/18/the-100-best-male-footballers-in-the-world-2018-no-s-100-71>.
 - [9] The Guardian, "The 100 best male footballers in the world 2019," *The Guardian*, Dec. 20, 2019. [Online]. Available: <https://www.theguardian.com/global/ng-interactive/2019/dec/17/the-100-best-male-footballers-in-the-world-2019>.
 - [10] M. Christenson, "The 100 best male footballers in the world 2020," *The Guardian*, Dec. 24, 2020. [Online]. Available: <https://www.theguardian.com/football/ng-interactive/2020/dec/21/the-100-best-male-footballers-in-the-world-2020>.
 - [11] The Guardian, "The 100 best male footballers in the world 2021," *The Guardian*, Dec. 24, 2021. [Online]. Available: <https://www.theguardian.com/football/ng-interactive/2021/dec/21/the-100-best-male-footballers-in-the-world-2021>.
 - [12] The Guardian, "The 100 best male footballers in the world 2022," *The Guardian*, Jan. 27, 2023. [Online]. Available: <https://www.theguardian.com/football/ng-interactive/2023/jan/24/the-100-best-male-footballers-in-the-world-2022>.
 - [13] M. Christenson, G. Blight, S. Clarke, S. Bloor, and P. Duncan, "The 100 best male footballers in the world 2023," *The Guardian*, Dec. 22, 2023. [Online]. Available: <https://www.theguardian.com/football/ng-interactive/2023/dec/19/the-100-best-male-footballers-in-the-world-2023>.
 - [14] M. Christenson, P. Duncan, S. Bloor, and G. Blight, "The 100 best male footballers in the world 2024," *The Guardian*, Dec. 20, 2024. [Online]. Available: <https://www.theguardian.com/football/ng-interactive/2024/dec/17/the-100-best-male-footballers-in-the-world-2024>.
 - [15] One Football, *Paper reveals list of the 100 best players of 2024/25, have your say*, <https://onefootball.com/en/news/paper-reveals-list-of-the-100-best-players-of-202425-have-your-say-41420267>, Jul. 24, 2025.
 - [16] S. Mulani, *StandardScaler in Python: Scikit-learn Guide & Examples*, <https://www.digitalocean.com/community/tutorials/standardscaler-function-in-python>, Aug. 3, 2022.
 - [17] GeeksforGeeks, *How Does the class_weight Parameter in Scikit-Learn Work?* <https://www.geeksforgeeks.org/machine-learning/how-does-the-classweight-parameter-in-scikit-learn-work/>, Jul. 23, 2025.
 - [18] S. Judge, "MLS, Saudi showing they are worth the switch," *Hindustan Times*, Aug. 12, 2025. [Online]. Available: <https://www.hindustantimes.com/sports/football/mls-saudi-showing-they-are-worth-the-switch-101755015388127.html>.
 - [19] scikit-learn, *RandomForestClassifier*, <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>.
 - [20] scikit-learn, *ExtraTreesClassifier*, <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.ExtraTreesClassifier.html>.
 - [21] scikit-learn, *RandomizedSearchCV*, https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.RandomizedSearchCV.html.
 - [22] Open AI, *ChatGPT*, <https://chatgpt.com/>.
 - [23] V. Favilla, *Regularization in Logistic Regression*, <https://medium.com/@vincefav/regularization-in-logistic-regression-14b50d7cc31>, May 29, 2023.
 - [24] scikit-learn, *LogisticRegression*, https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html.
 - [25] scikit-learn, *1.4. Support Vector Machines*, https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html.
 - [26] M. Ben Fraj, *In Depth: Parameter tuning for Random Forest*, <https://medium.com/all-things-ai/in-depth-parameter-tuning-for-random-forest-d67bb7e920d>, Dec. 21, 2017.
 - [27] GeeksforGeeks, *Bootstrap Method*, <https://www.geeksforgeeks.org/maths/bootstrap-method/>, Jul. 23, 2025.
 - [28] E. Clower, *Understanding Cross-Validation*, <https://www.aptech.com/blog/understanding-cross-validation/>, Mar. 20, 2025.
 - [29] A. Singh, *Cross-Validation Techniques*, <https://medium.com/geekculture/cross-validation-techniques-33d389897878>, Aug. 30, 2021.
 - [30] GeeksforGeeks, *Understanding Feature Importance in Logistic Regression Models*, <https://www.geeksforgeeks.org/machine-learning/understanding-feature-importance-in-logistic-regression-models/>, Jul. 23, 2025.

- [31] GeeksforGeeks, *Feature Importance with Random Forests*, <https://www.geeksforgeeks.org/machine-learning/feature-importance-with-random-forests/>, Nov. 11, 2025.
- [32] S. Galli, *Understanding Permutation Feature Importance for Model Interpretation*, <https://www.blog.trainindata.com/permutation-feature-importance/>, Jan. 11, 2024.
- [33] V. Satheesh, *Hyper Parameter Tuning (GridSearchCV Vs RandomizedSearchCV)*, <https://medium.com/analytics-vidhya/hyper-parameter-tuning-gridsearchcv-vs-randomizedsearchcv-499862e3ca5>, Dec. 22, 2020.
- [34] M. Matthews, “Every Ballon d’Or Winner by Position,” *GiveMeSport*, May 26, 2025. [Online]. Available: <https://www.givemesport.com/every-ballon-dor-winner-by-position-striker-winger-messi/>.
- [35] B. Miller, “Who is Lamine Yamal? Contract, playing style, La Liga records broken by Barcelona and Spain sensation,” *The Sporting News*, Jul. 9, 2024. [Online]. Available: <https://www.sportingnews.com/us/soccer/news/lamine-yamal-barcelona-la-liga-spain-records-playing-style/723f032776a836c20966f8ac>.
- [36] Wikipedia, *Learning to rank*, https://en.wikipedia.org/wiki/Learning_to_rank, Apr. 9, 2026.
- [37] M. Christenson, P. Duncan, S. Bloor, and G. Blight, “The 100 best male footballers in the world 2025,” *The Guardian*, Dec. 19, 2025. [Online]. Available: <https://www.theguardian.com/football/ng-interactive/2025/dec/16/the-100-best-male-footballers-in-the-world-2025>.

Appendix A




Figure 1: Test Model Percentages Bar Graph




Figure 2: Full Percentages Bar Graph

Predicted Ranking Matches Marca Ranking
Correctly Predicted in Marca Top 100

	Marca Rankings	LR Rankings	RBF Rankings	Sigmoid Rankings	Poly Rankings	RF Rankings
1	Lamine Yamal	Kylian Mbappé	Raphinha	Joshua Kimmich	Kylian Mbappé	Ousmane Dembélé
2	Ousmane Dembélé	Mohamed Salah	Omar Marmoush	Marcos Alonso	Mohamed Salah	Kylian Mbappé
3	Vitinha	Raphinha	Bradley Barcola	Vitinha	Raphinha	Cole Palmer
4	Raphinha	Ousmane Dembélé	Kylian Mbappé	Joško Gvardiol	Lamine Yamal	Michael Olise
5	Kylian Mbappé	Bradley Barcola	Cole Palmer	Pedri	Omar Marmoush	Alexander Isak
6	Pedri	Joshua Kimmich	Alexander Isak	Hakan Çalhanoğlu	Bradley Barcola	Bradley Barcola
7	Mohamed Salah	Lamine Yamal	Ousmane Dembélé	Achraf Hakimi	Harry Kane	Raphinha
8	Achraf Hakimi	Michael Olise	Mohamed Salah	Serge Gnabry	Michael Olise	Mohamed Salah
9	Nuno Mendes	Omar Marmoush	Michael Olise	Nadiem Amiri	Cole Palmer	Mason Greenwood
10	Gianluigi Donnarumma	Harry Kane	Hugo Ekitike	Federico Valverde	Mateo Retegui	Omar Marmoush
11	Désiré Doué	Mateo Retegui	Mason Greenwood	Ferrán Torres	Joshua Kimmich	Lamine Yamal
12	Viktor Gyökeres	Alexander Isak	Mateo Retegui	Harvey Barnes	Ousmane Dembélé	Hugo Ekitike
13	Harry Kane	Hugo Ekitike	Robert Lewandowski	Himad Abdelli	Vinicius Júnior	Harry Kane
14	Cole Palmer	Pierre Højbjerg	Harry Kane	Marco Asensio	Alexander Isak	Joshua Kimmich
15	Lautaro Martínez	Robert Lewandowski	Florian Wirtz	Adrien Rabiot	Florian Wirtz	Erling Haaland
16	Julián Álvarez	Luis Díaz	Luis Díaz	Martin Ødegaard	Hugo Ekitike	Luis Díaz
17	Fabián Ruiz Peña	Pedri	Vinicius Júnior	Dani Olmo	Luis Díaz	Achraf Hakimi
18	Robert Lewandowski	Florian Wirtz	Joshua Kimmich	Luis Henrique	Mason Greenwood	Mateo Retegui
19	Scott McTominay	Mason Greenwood	Lamine Yamal	James Maddison	Pedri	Ollie Watkins
20	Declan Rice	Achraf Hakimi	Erling Haaland	Alassane Pléa	Robert Lewandowski	Granit Xhaka
21	Federico Valverde	Vitinha	Serhou Guirassy	Deniz Undav	Rayan Cherki	Julián Álvarez
22	Isco	Désiré Doué	Rafael Leão	Cody Gakpo	Désiré Doué	Vinicius Júnior
23	João Neves	Serhou Guirassy	Matheus Cunha	Borja Iglesias	Maghnes Aklouché	Rafael Leão
24	Florian Wirtz	Granit Xhaka	Pedri	Enzo Fernández	Serhou Guirassy	Florian Wirtz
25	Vinicius Júnior	Vinicius Júnior	Khvicha Kvaratskhelia	Eliesse Ben Seghir	Granit Xhaka	Serhou Guirassy
26	Denzel Dumfries	Khvicha Kvaratskhelia	Désiré Doué	Álex Grimaldo	Jude Bellingham	Matheus Cunha
27	Virgil van Dijk	Pau Cubarsí	Rayan Cherki	Brennan Johnson	Rafael Leão	Pierre Højbjerg
28	Jamal Musiala	Cole Palmer	Bryan Mbeumo	Youri Tielemans	Khvicha Kvaratskhelia	Pedri
29	Alessandro Bastoni	João Neves	Bruno Fernandes	Bruno Guimarães	Matheus Cunha	Ademola Lookman
30	Cristiano Ronaldo	Jude Bellingham	Maghnes Aklouché	Breel Embolo	Ademola Lookman	Maghnes Aklouché
31	Yann Sommer	Joško Gvardiol	Granit Xhaka	Tijani Reijnders	Marcos Alonso	Désiré Doué
32	Michael Olise	Bruno Fernandes	Ademola Lookman	Kingsley Coman	Achraf Hakimi	Lewis Dunk

Figure 3: Rankings Part 1

33	Alexander Isak	Leroy Sané	Martin Ødegaard	Kevin Schade	Martin Ødegaard	Khvicha Kvaratskhelia
34	Nico Williams	Rayan Cherki	Jamal Musiala	Adam Hložek	Bruno Fernandes	Christian Pulisic
35	Jules Koundé	Angelo Stiller	Alex Iwobi	Enzo Millot	Jamal Musiala	Nicolò Barella
36	Pau Cubarsí	Julián Álvarez	Achraf Hakimi	Jørgen Strand Larsen	Bukayo Saka	Rayan Cherki
37	Bukayo Saka	Kim Min-jae	Jude Bellingham	Emanuel Emegha	Son Heung-min	Iñigo Martínez
38	Antony	William Saliba	Julián Álvarez	Dominic Solanke	Pierre Højbjerg	Bruno Fernandes
39	Willian Pacho	Yoane Wissa	Angelo Stiller	Rayan Aït-Nouri	Iñigo Martínez	Éderson
40	Erling Haaland	Leonardo Balerdi	Son Heung-min	Corentin Tolisso	Vinha	Federico Valverde
41	Jude Bellingham	Virgil van Dijk	Oihan Sancet	Mitchell Weiser	Luis Henrique	Antoine Semenyo
42	Luis Díaz	Mateo Kovačić	Nicolás Paz	Gabriel Martinelli	Federico Valverde	Jude Bellingham
43	Dean Huijsen	Fabián Ruiz Peña	Moise Kean	Karim Adeyemi	Luka Modrić	Moise Kean
44	William Saliba	Matheus Cunha	Luka Modrić	Abde Ezzalzouli	Trent Alexander-Arnold	Angelo Stiller
45	Joan García	Maghnes Aklouch	Gaëtan Perrin	Pierre Højbjerg	Kaoru Mitoma	Leroy Sané
46	Rayan Cherki	Federico Valverde	Leroy Sané	Randal Kolo Muani	Angelo Stiller	Antonio Rüdiger
47	Thibaut Courtois	Son Heung-min	Trent Alexander-Arnold	Scott McTominay	Julián Álvarez	Nicolás Paz
48	Marcus Thuram	Iñigo Martínez	Marcos Alonso	Evanilson	Pau Cubarsí	Fabián Ruiz Peña
49	Lionel Messi	Vincenzo Grifo	Pierre Højbjerg	Iliman Ndiaye	Joško Gvardiol	Martin Ødegaard
50	Luka Modrić	Ayoze Pérez	Ollie Watkins	Mika Biereth	Ollie Watkins	Bruno Guimarães
51	Romelu Lukaku	Alexander Sørloth	Arnaud Kalimuendo	Alejandro Garnacho	Amad Diallo	Virgil van Dijk
52	Tijjani Reijnders	Rafael Leão	Jonathan Burkardt	Theo Hernández	Leroy Sané	Kim Min-jae
53	Jan Oblak	Erling Haaland	Kaoru Mitoma	Rodrygo	Fabián Ruiz Peña	Luka Modrić
54	Mateo Retegui	Luka Modrić	Yoane Wissa	Boulaye Dia	Vincenzo Grifo	Alex Iwobi
55	Alex Baena	Ademola Lookman	Christian Pulisic	Danny Welbeck	Ayoze Pérez	Bernardo Silva
56	Bruno Fernandes	Amir Rrahmani	Bukayo Saka	Robin Hack	James Maddison	Alessandro Bastoni
57	Iñigo Martínez	Alessandro Bastoni	Fabián Ruiz Peña	Amad Diallo	Kingsley Coman	Yoane Wissa
58	Ante Budimir	James Maddison	Federico Valverde	Joshua King	Valentín Castellanos	Nicolas Jackson
59	Ayoze Pérez	Jonathan Burkardt	Jarrod Bowen	Lucas Stassin	Virgil van Dijk	Andrej Kramarić
60	Rodrigo De Paul	Andrey Santos	Gonçalo Ramos	Romano Schmid	Xavi Simons	Rúben Dias
61	Moisés Caicedo	Bafodé Diakité	Iñigo Martínez	Che Adams	Erling Haaland	Jarrod Bowen
62	Joshua Kimmich	Alex Iwobi	Lee Kang-in	Patrick Cutrone	Leandro Trossard	Jonathan Tah
63	Karim Benzema	Nico Schlotterbeck	Antoine Semenyo	Sebastian Nanasi	Alex Iwobi	Pau Cubarsí
64	Serhou Guirassy	Amine Gouiri	Tijjani Reijnders	Santiago Castro	Zuriko Davitashvili	Joško Gvardiol
65	Mason Greenwood	Nadiem Amiri	Bruno Guimarães	Antoine Griezmann	Nikola Krstović	Sávio
66	Ryan Gravenberch	Álex Grimaldo	Ante Budimir	Dennis Man	Dominik Szoboszlai	Jonathan Burkardt

Figure 4: Rankings Part 2

67	Trent Alexander-Arnold	Lucas Beraldo	Joško Gvardiol	Giovani Lo Celso	Daley Blind	Son Heung-min
68	Nicolás Paz	Jonathan David	Pau Cubarsí	Oihan Sancet	Georges Mikautadze	Zuriko Davitashvili
69	Alejandro Balde	Adrien Rabiot	Nico Williams	Keito Nakamura	Federico Dimarco	Jean-Philippe Mateta
70	Omar Marmoush	Jamal Musiala	João Neves	Hamed Junior Traorè	Alessandro Bastoni	Hakan Çalhanoğlu
71	Victor Osimhen	Hakan Çalhanoğlu	Charles De Ketelaere	Ludovic Ajorque	Jubal	İlkay Gündoğan
72	Francesco Acerbi	Rúben Dias	Amine Gouiri	Lee Kang-in	Cody Gakpo	Bukayo Saka
73	Matheus Cunha	Corentin Tolisso	Alexander Sørloth	Maximilian Beier	Nuno Tavares	Álex Grimaldo
74	Oihan Sancet	Martin Ødegaard	James Maddison	Federico Dimarco	Tijjani Reijnders	Amir Rrahmani
75	Álvaro Carreras	Marcos Alonso	Xavi Simons	Zakaria Aboukhilal	Amine Gouiri	James Maddison
76	Francisco Trincão	Ollie Watkins	Ayoze Pérez	Diego López	Diego Moreira	Thilo Kehrer
77	Fábio Silva	Dayot Upamecano	Jonathan David	Fabián Ruiz Peña	Nicolás Paz	Eberechi Eze
78	Moise Kean	Serge Gnabry	Eberechi Eze	Leroy Sané	Jonathan Burkardt	Manuel Akanji
79	Pedro Porro	Kaoru Mitoma	Dani Olmo	Kai Havertz	Oihan Sancet	Mattéo Guendouzi
80	Riyad Mahrez	Antoine Griezmann	Álex Grimaldo	Xavi Simons	Lee Kang-in	Koke
81	Ivan Perišić	Marcus Thuram	Enzo Fernández	Kevin De Bruyne	Marten de Roon	Marquinhos
82	Bruno Guimarães	Luis Henrique	Arda Güler	Patrik Schick	Marcus Thuram	Amine Gouiri
83	Vangelis Pavlidis	Emanuel Emegha	Kingsley Coman	Paulo Dybala	Kim Min-jae	Xavi Simons
84	Hugo Ekitike	Tijjani Reijnders	Alessandro Bastoni	Ermedin Demirović	Rayan Aït-Nouri	Enzo Fernández
85	Murillo	Rayan Aït-Nouri	Vinha	Dango Ouattara	Sávio	Ibrahima Konaté
86	Xherdan Shaqiri	Bryan Mbeumo	Jubal	Granit Xhaka	Emre Can	Nico Schlotterbeck
87	Mika Biereth	Levi Colwill	Luis Henrique	Mattia Zaccagni	João Neves	Amad Diallo
88	Andrey Santos	Daley Blind	Fer López	Loïs Openda	Charles De Ketelaere	Bryan Mbeumo
89	Óscar Mingueza	Lee Kang-in	Marten de Roon	Jorge de Frutos	Jonathan Tah	Michael Amir Murillo
90	Amir Rrahmani	Enzo Fernández	Fermin López	Paul Nebel	Benjamin Šeško	Mateo Kovačić
91	Eberechi Eze	Bruno Guimarães	Georges Mikautadze	Marcus Thuram	Serge Gnabry	Marten de Roon
92	Bryan Mbeumo	Moise Kean	Virgil van Dijk	Thierno Barry	Alexander Sørloth	Dominik Szoboszlai
93	Sem Steijn	Jules Koundé	Mika Biereth	Charles De Ketelaere	Tom Bischof	Levi Colwill
94	Mile Svilar	Lautaro Martínez	Cody Gakpo	Julian Brandt	Gonçalo Ramos	Vincenzo Grifo
95	David de Gea	Marquinhos	Rodrigo De Paul	Gustav Isaksen	Nico Williams	Robert Lewandowski
96	Mark Flekken	Patrik Schick	Dominik Szoboszlai	Jamie Gittens	Nico Schlotterbeck	Romelu Lukaku
97	Amad Diallo	Éderson	Nico Schlotterbeck	Kenny Lala	Kenan Yıldız	Tijjani Reijnders
98	Kasper Waarst Høgh	Jonathan Tah	Gorka Rivera Alonso	Takumi Minamino	Tim Kleindienst	Federico Dimarco
99	Daizen Maeda	Gonçalo Ramos	Kim Min-jae	Shuto Machino	Romelu Lukaku	Thomas Partey
100	Oliver Antman	Kingsley Coman	Federico Dimarco	Angelo Stiller	Marco Friedl	Jules Koundé

Figure 5: Rankings Part 3